

C++ 14 Language Features

Contents

a)	Core Language Features.....	2
1.	Digit Separators	2
2.	Binary Literals	2
3.	Generic Lambdas	3
4.	Generalized Lambda Captures	4
5.	Return Type Deduction for Functions.....	5
6.	The <code>decltype(auto)</code>	6
7.	Relaxed <code>constexpr</code> Restrictions	7
8.	Variable Templates.....	9
9.	Deprecated Attribute	9
10.	Sized Deallocation.....	10
b)	Class-Related Features	12
11.	Default Member Initializers for Bit-Fields.....	12
12.	Relaxed Requirements for <code>constexpr</code> Member Functions.....	12
13.	Inheriting Constructors	13
14.	Aggregate Initialization with Member Initializers	14
c)	Standard Library Features.....	16
15.	Standard User-Defined Literals	16
16.	The <code>std::make_unique</code>	16
17.	The <code>std::exchange</code>	17
18.	Heterogeneous Lookup in Associative Containers.....	18
19.	The <code>std::quoted</code>	19
20.	The <code>std::integer_sequence</code>	20

a) Core Language Features

1. Digit Separators

- In **C++14**, **digit separators** (') were introduced to improve the readability of numeric literals.
- The **digit separator** is a **single quote** (') that you can place anywhere inside a numeric literal to make large numbers easier to read. It has **no effect on the value**; it is purely for formatting.
- Why are Digit Separators Useful?
 - When dealing with large numbers, it can be difficult to read and understand them at a glance. The digit separator allows programmers to group digits in a meaningful way, similar to how commas are used in everyday numbers.
 - For example:
 - 1000000 is harder to read.
 - 1'000'000 is easier to read.

```
using namespace std;
int main() {
    int oneMillion = 1'000'000;      // 1 million
    int tenLakh = 10'00'000;        // 1 million = 10 lakh
    cout << "1 million: " << oneMillion << endl;
    cout << "10 lakh: " << tenLakh << endl;
    return 0;
}
```

2. Binary Literals

- A **binary literal** is a way to specify numbers in base-2 (binary) directly in C++ code using the **0b** or **0B** prefix.
- This allows programmers to define numbers using only **0s** and **1s**, which is useful for low-level programming, bitwise operations, and hardware programming.
- Example 1: A Simple Bitmask

```
const int FEATURE_A = 0b0000'0001; // Bit 0
const int FEATURE_B = 0b0000'0010; // Bit 1
const int FEATURE_C = 0b0000'0100; // Bit 2
```

```

int main() {
    int binaryNumber = 0b101010; // 42 in decimal
    cout << "Binary 0b101010 is: " << binaryNumber << endl;
    int anotherBinary = 0b1111; // 15 in decimal
    cout << "Binary 0b1111 is: " << anotherBinary << endl;
    int largeBinary = 0b1101'0101'1010'0110; // Grouped for better readability
    cout << "Binary 0b1101'0101'1010'0110 is: " << largeBinary << endl;
}

```

3. Generic Lambdas

- A **lambda function** is an **anonymous function** (a function without a name) that you can define inline.
- In **C++14**, you can use **auto** as a lambda parameter type to make the lambda **type-generic** (like a function template).

```

int main() {
    auto add = [] (auto a, auto b) {
        return a + b;
    }; // Works with any type
    cout << add(3, 4) << endl; // Works with int (3 + 4)
    cout << add(3.5, 2.1) << endl; // Works with double (3.5 + 2.1)
    cout << add(string("Hi "), "C++!") << endl; // Works with strings
}

```

- Why Use Generic Lambdas?
 - **More Flexible:** Works with any data type (**int**, **double**, **string**, etc.).
 - **Avoids Code Duplication:** No need to write multiple versions of the same function for different types.
 - **Simplifies Code:** You don't have to write explicit template functions.

Comparison: C++11 vs. C++14

Feature	C++11 (Fixed Type)	C++14 (Generic)
Syntax	<code>[](int a, int b) { return a + b; }</code>	<code>[](auto a, auto b) { return a + b; }</code>
Works with int ?	✓	✓
Works with double ?	✗ (Would need another lambda)	✓
Works with string ?	✗	✓

4. Generalized Lambda Captures

- What Are Lambdas?
 - Lambdas, or lambda expressions, are a way to define small, unnamed functions directly in your code. They can be very useful for quick, short-lived tasks, especially when working with algorithms and callbacks.
- What Are Lambda Captures?
 - Lambda captures allow lambdas to **capture and use variables from their enclosing scope** (the scope in which they are defined). This can be useful when you want the lambda to have access to local variables.
- Generalized Lambda Captures in C++14
 - Before C++14, lambdas could **only capture variables by value ([=]) or by reference ([&]).**
 - C++14 introduced **generalized lambda captures**, allowing us to **move objects** and **initialize new variables inside the capture list.**
- Example:
 - Before C++14, capturing **move-only** objects (like `std::unique_ptr`) was not possible.
 - Before C++14, Capture expressions (e.g., `x + 5`) were not possible.

```
int main() {
    auto ptr = std::make_unique<int>(42);

    // Move ptr into lambda using generalized capture
    auto lambda1 = [uniquePtr = std::move(ptr)]() {
        std::cout << *uniquePtr << "\n";
    };
    lambda1(); // Output: 42

    auto x = 10;
    auto lambda2 = [y = x + 5]() { // Capturing an expression (x + 5)
        std::cout << y << "\n";
    };
    lambda2();
}
```

5. Return Type Deduction for Functions

- In **C++14, Return Type Deduction for Functions** was introduced to make function definitions more flexible by allowing the compiler to deduce the return type automatically.
- What's Return Type Deduction?
 - In C++, every function has a return type – the type of value the function gives back when it's finished. Before C++14, you had to write this return type explicitly in the function's declaration.
 - Example:

```
// Explicit return type: int
int add(int x, int y) {
    return x + y;
}

// Explicit return type: double
double divide(double x, double y) {
    return x / y;
}
```

- In **C++14**, you can **omit the return type**, and the compiler will automatically deduce it from the **return** statement.

```
// Compiler deduces return type
auto add(int a, int b) {
    return a + b;
}
```

C++11 vs. C++14 Return Type Deduction

Feature	C++11	C++14
Function Return Type	Must be explicitly declared	Can be deduced with auto
Syntax	<code>int add(int a, int b){ return a + b; }</code>	<code>auto add(int a, int b) { return a + b; }</code>
Works with Iterators?	Must specify complex type	Deduces iterator type automatically
Template Functions?	Requires explicit return type	Deduces return type automatically

6. The `decltype(auto)`

- In C++14, `decltype(auto)` was introduced to improve **return type deduction**, making it more flexible than just using `auto`.
- What is `decltype`?
 - `decltype` is an operator that lets you get the *type* of an expression.
 - Example:

```
int x = 10;
decltype(x) y = 20; // y is of type int (same as x)
```

- The Problem: Type Decay and References
 - Sometimes, `decltype` doesn't give you *exactly* the type you might expect, especially when dealing with functions that return references or when dealing with *type decay*.

```
int& get_value() {
    static int value = 42;
    return value;
}

int main() {
    // Type of ref1 is int
    auto ref1 = get_value();
    ref1 = 50;
    cout << "New Value: " << get_value();
}
// Output: New Value: 42
```

- What is `decltype(auto)`?
 - `decltype(auto)` allows the compiler to **deduce** the exact type, just like `decltype(expr)` which evaluates the type of the expression.
 - It preserves references and `const` qualifiers, unlike `auto`, which may strip them away.
 - It is useful for `return` type deduction when you want to maintain the original type precisely.

- Example:

```

int& get_value() {
    static int value = 42;
    return value;
}

int main() {
    // Type of ref1 is int&
    decltype(auto) ref1 = get_value();
    ref1 = 50;
    cout << "New Value: " << get_value();
}
// Output: New Value: 50

```

Comparison Table

Feature	auto	decltype	decltype(auto)
Deduces type?	✓ Yes	✗ No (only evaluates type)	✓ Yes
Removes reference?	✓ Yes	✗ No	✗ No
Removes const?	✓ Yes	✗ No	✗ No
Used for function return type?	✓ Yes	✗ No (used for type evaluation)	✓ Yes
Supports complex expressions?	✓ Yes	✓ Yes	✓ Yes

7. Relaxed `constexpr` Restrictions

- The `constexpr` is a keyword that you can use to **declare variables or functions** that *can be evaluated at compile time*.

The Problem: Restrictions in C++11

- In C++11, `constexpr` functions had very strict limitations.
 - They **could not have loops** (`for`, `while`, etc.).
 - They **could not use local variables**.
 - They **could not use if statements**.
 - Every `constexpr` function had to be one `return` statement.

The Solution: Relaxed Restrictions in C++14

- C++14 relaxed these rules, making **constexpr** functions more flexible.
- Now, **constexpr** functions can have:
 - **Loops** (**for**, **while**, **do-while**)
 - **Multiple statements** (instead of one **return**)
 - **Local variables** (with **auto**)
 - **Conditional statements** (**if**, **switch**, etc.)
- Example:

```
// Only single return allowed in C++11
constexpr int factorial(int n) {
    return (n <= 1) ? 1 : (n * factorial(n - 1));
}

constexpr int factorial(int n) {
    int result = 1;
    // Loops are allowed in C++14
    for (int i = 2; i <= n; ++i) {
        result *= i;
    }
    return result;
}
```

Summary: What Changed in C++14?

Feature	C++11	C++14
Multiple statements	✗ No	✓ Yes
Loops (for , while)	✗ No	✓ Yes
Local variables	✗ No	✓ Yes
Conditional statements (if , switch)	✗ No	✓ Yes
Modifying variables	✗ No	✓ Yes

8. Variable Templates

- Templates allow us to write generic code that works for multiple types.
- Before **C++14**, templates were used only for **functions** and **classes**.
- In **C++14**, we can now create **variable templates**.

What is a Variable Template?

- A variable template allows defining a **variable with a generic type**, just like function templates.
- Example:

```
// Define a variable template
template<typename T>
constexpr T pi = T(3.1415926535897932385);

int main() {
    cout << pi<int> << "\n";
    cout << fixed << setprecision(7);
    cout << pi<float> << "\n";
    cout << fixed << setprecision(15);
    cout << pi<double> << "\n";
}
/*
3
3.1415927
3.141592653589793
*/
```

- The same **pi** variable works for different types (**int**, **float**, **double**).
- The **constexpr** ensures compile-time computation.
- No need to create multiple constants for each type.

9. Deprecated Attribute

- The **[[deprecated]] attribute** was introduced in **C++14** to mark **functions**, **classes**, **variables**, or **enumerations** as outdated or obsolete. This helps warn developers when they use old code that might be removed in the future.

Why Do We Need `[[deprecated]]`?

- Imagine you're working on a large C++ project, and some functions or classes are no longer recommended because:
 - They are inefficient (e.g., replaced with better alternatives).
 - They have security issues.
 - They will be removed in future versions.
- Instead of removing the function immediately, you mark it as deprecated so that existing code still works, but compilers show a warning when someone tries to use it.

```
// Marking a function as deprecated with a message
[[deprecated("Use newFunction() instead.")]]
void oldFunction() {
    cout << "This function is deprecated!\n";
}
// New recommended function
void newFunction() {
    cout << "This is the new function!\n";
}
struct [[deprecated("Use NewClass instead.")]] OldClass {
    int value;
};
struct NewClass {
    int value;
};
int main() {
    OldClass obj; // Warning: "Use NewClass instead."
    oldFunction(); // Warning with message: "Use newFunction() instead."
    newFunction(); // No warning
}
```

10. Sized Deallocation

- In **C++14**, the memory deallocation functions (`operator delete` and `operator delete[]`) can optionally receive the **size of the object** being deleted. This is known as **sized deallocation**.

Why is Sized Deallocation Useful?

- Optimization for Custom Allocators:

- Some custom memory allocators need to know the exact size of the object when freeing memory.
- Before C++14, `operator delete` had to figure out the size by other means, which could be inefficient.
- Potential Performance Improvements:
 - If the deallocation function knows the size upfront, it can free memory faster.
 - This helps in memory pools or custom allocators that manage different object sizes.

How Does It Work?

- In C++14, the compiler can pass the `size` of the object to `operator delete`, allowing optimized memory deallocation.
- C++11 and Before:

```
void operator delete(void* ptr) noexcept;
void operator delete[](void* ptr) noexcept;
```
- C++14 (Sized Deallocation Added):

```
void operator delete(void* ptr, std::size_t size) noexcept;
void operator delete[](void* ptr, std::size_t size) noexcept;
```

b) Class-Related Features

11. Default Member Initializers for Bit-Fields

- In C++14, **bit-fields** (integer members with a specified number of bits) can now have **default values** directly in the class or struct definition.

```
struct Flags {  
    unsigned int a : 1 {1}; // Default: 1  
    unsigned int b : 2 {0}; // Default: 0  
    unsigned int c : 3 {5}; // Default: 5  
};  
  
int main() {  
    Flags f; // Uses default values  
    std::cout << f.a << " " << f.b  
        << " " << f.c << "\n"; // Output: 1 0 5  
}
```

12. Relaxed Requirements for constexpr Member Functions

- What is **constexpr**?
 - The **constexpr** keyword is used to define functions and variables that can be evaluated at compile-time. This can help optimize your code by performing calculations during compilation rather than at runtime.
- What are Member Functions?
 - Member functions are functions that belong to a **class**. They can access and modify the data members of the **class**.

Changes in C++14

- Before C++14, **constexpr** member functions were implicitly **const**. This means that they couldn't modify the object's state (i.e., they couldn't change the values of data members).
- In C++14, the requirement that **constexpr** member functions be **const** was relaxed, allowing **constexpr** member functions to modify the object's state if necessary.
- The **constexpr** keyword ensures that the member functions can be evaluated at compile time if the context allows it, but it doesn't prevent them from being used at runtime.

```

// A class with constexpr member functions
class Counter {
private:
    int count;
public:
    // Constructor to initialize count
    constexpr Counter(int initialCount) : count(initialCount) {}
    // constexpr member function to increment the count
    constexpr Counter increment() const {
        return Counter(count + 1);
    }
    // constexpr member function to get the count
    constexpr int getCount() const {
        return count;
    }
};
int main() {
    // Create a Counter object using a constexpr expression
    constexpr Counter counter = Counter(0);
    counter.increment().increment();
    // Get the count at compile time
    constexpr int result = counter.getCount();
    cout << "Count1: " << result << endl;

    cout << "Enter a value: ";
    int i; cin >> i;
    Counter counter2 = Counter(i);
    counter2.increment().increment();
    auto result1 = counter2.getCount();
    cout << "Count2: " << result1 << endl;

    return 0;
}

```

13. Inheriting Constructors

- In C++14, the ability to **inherit constructors** from a **base class** was introduced, making it easier to write **classes** that need to inherit constructors from their base **classes**.

What is Constructor Inheritance?

- In C++11, when you create a derived **class**, you had to explicitly define constructors for the derived **class**, even if the constructor was just calling the base **class** constructor. This could lead to redundant code.

- With constructor inheritance (introduced in C++11) and improved in C++14, the derived **class** can inherit the constructor(s) from the base **class**, reducing redundancy.

```

class Base {
public:
    Base(int x) {
        cout << "Base constructor called with value: "
        << x << endl;
    }
};

class Derived : public Base {
public:
    // In C++11, we used to write a constructor in Derived
    // class that calls the constructor of the Base class.
    //Derived(int x) : Base(x) {
    //    cout << "Derived constructor called with value: "
    //    << x << endl;
    //}

    // C++14: Inherit constructors from Base class
    using Base::Base;
};

int main() {
    Derived obj(10);
    return 0;
}

```

14. Aggregate Initialization with Member Initializers

- In C++14, aggregate **classes** (simple **structures** or **classes** without constructors) can now **use default member initializers** while still allowing aggregate initialization.

What is Aggregate Initialization?

- Aggregate initialization is a way to initialize **structs** or **classes** **using braces {} without defining a constructor.**

```

struct Point {
    int x = 10; // Default member initializer
    int y = 20; // Default member initializer
};

```

```
int main() {
    // This was NOT allowed in C++11:
    // ERROR: Default member initializers prevent aggregate initialization.
    // Point p = {5, 6};

    Point p;          // Works (uses default values: x=10, y=20)
    std::cout << p.x << ", " << p.y << "\n";

    // C++14 allows aggregate initialization with default member initializers
    Point p1;          // Uses default values: x = 10, y = 20
    Point p2 = { 5 };    // x = 5, y = 20 (default is used for missing values)
    Point p3 = { 3, 4 }; // x = 3, y = 4 (all values provided)
}
```

c) Standard Library Features

15. Standard User-Defined Literals

- In C++14, **standard user-defined literals** were introduced, allowing developers to create more readable and expressive code with **custom literals** for common types like `std::string`, `std::chrono`, and `std::complex`.

What Were the Limitations Before C++14?

- Before C++14, it was possible to define custom literals, but there were no built-in support for literals of common types like `std::string`, `std::chrono`, or `std::complex`. This meant you had to define your own literal operators for these types.

```
using namespace std;
using namespace std::chrono;
int main() {
    std::string s = "Hello World!"s;
    auto duration = 5s;
    std::complex<double> c = 3.0 + 4.0i;

    cout << "String: " << s << endl;
    cout << "Duration: " << duration.count()
        << " seconds" << endl;
    cout << "Complex: " << c << endl;
    return 0;
}
```

16. The `std::make_unique`

- In C++14, the `std::make_unique` function was introduced as a **safe, efficient, and convenient** way to create `std::unique_ptr` objects.

```
int main() {
    // C++11: But using new manually is error-prone and less readable.
    std::unique_ptr<int> ptr1(new int(42)); // Manually allocating memory
    std::cout << *ptr1 << std::endl; // Output: 42

    // C++14: std::make_unique() is a better way.
    auto ptr2 = std::make_unique<int>(42); // Creates a unique_ptr<int>
    std::cout << *ptr2 << std::endl; // Output: 42
    return 0;
}
```

Summary

- `std::make_unique` (introduced in C++14) is a safer, cleaner, and more efficient way to create `std::unique_ptr`.
- It eliminates the need for `new`, preventing memory leaks and improving exception safety.
- Supports single objects (`std::make_unique<int>(42)`) and arrays (`std::make_unique<int[]>(5)`).
- Use it whenever possible, except when a custom `deleter` is required.

Quick Comparison: `new` vs. `std::make_unique`

Feature	<code>new</code> with <code>std::unique_ptr</code>	<code>std::make_unique</code>
Memory Safety	✗ Manual <code>new/delete</code> needed	✓ No need for <code>new</code> or <code>delete</code>
Exception Safety	✗ Risk of leaks if an exception occurs	✓ Ensures safe memory allocation
Code Readability	✗ Verbose and error-prone	✓ Concise and clean
Array Support	✓ Supports <code>new[]</code> but requires <code>delete[]</code>	✓ Supports arrays automatically

17. The `std::exchange`

- The `std::exchange` is a utility function introduced in C++14 that:
 - Replaces the value of an object with a new value.
 - Returns the old value of the object before the replacement.

```
#include <utility> // exchange

int main() {
    int x = 10;
    // x becomes 20, and old_value stores 10
    int old_value = exchange(x, 20);

    cout << "Old Value: " << old_value << "\n"; // Output: 10
    cout << "New Value of x: " << x << "\n"; // Output: 20

    return 0;
}
```

18. Heterogeneous Lookup in Associative Containers

What is Heterogeneous Lookup?

- Before C++14, when searching for a key in associative containers (`std::set`, `std::map`, `std::unordered_map`, etc.), you had to use exactly the same key type.
- C++14 introduced heterogeneous lookup, which allows searching using different but compatible types (e.g., looking up a `std::string` key using a `const char*`).

Why is This Useful?

- **Efficiency:** Avoids unnecessary conversions when searching.
- **Flexibility:** You can search using a different but compatible type.
- **Performance:** Reduces object creation overhead.

```
int main() {
    std::set<std::string> mySet = { "apple", "banana", "cherry" };

    // Searching using const char* (does NOT work in C++11)
    // auto it = mySet.find("banana"); // ERROR: Must convert to std::string
    auto it1 = mySet.find(std::string("banana")); // Works but inefficient

    // C++14: Searching using const char*
    auto it2 = mySet.find("banana");
    if (it2 != mySet.end())
        std::cout << "Found: " << *it2 << "\n";
}
```

How Does This Work?

- C++14 Adds Transparent Comparators (`std::less<>`)
- C++14 allows `std::set`, `std::map`, etc., to use **transparent comparators**, such as `std::less<>` instead of `std::less<T>`.

But what are transparent comparators?

- A **transparent comparator** is a special type of comparator that allows **heterogeneous lookup**—meaning you can **search** for elements in a `std::set` or `std::map` using different but compatible types **without unnecessary conversions**.

Summary

Feature	<code>std::less<T></code>	<code>std::less<>(Transparent)</code>
Requires exact type match?	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No
Supports <code>const char*</code> vs <code>std::string</code> ?	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes
Unnecessary object creation?	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No
Performance overhead?	<input checked="" type="checkbox"/> High	<input type="checkbox"/> Low
Introduced in	C++98	C++14

19. The `std::quoted`

- The `std::quoted` is a C++14 feature that simplifies handling quoted strings when using input (`>>`) and output (`<<`) streams.
- It automatically adds or removes quotes ("") when reading or writing strings.
- It escapes special characters like " and \ properly.

Why Do We Need `std::quoted`?

- Before C++14, managing quoted strings manually was difficult.

- Printing strings with quotes manually:

```
std::cout << "\" " << myString << "\" ;
```

- Example:

```
#include <iomanip> // Required for std::quoted

int main() {
    std::string name = "John Doe";

    // Writing to output stream (with quotes)
    std::cout << std::quoted(name) << "\n";
}
// Output: "John Doe"
```

Summary: Why Use `std::quoted`?

Feature	Without <code>std::quoted</code>	With <code>std::quoted</code>
Adding Quotes on Output	✗ Manual (" + str + ")	✓ Automatic (<code>std::quoted(str)</code>)
Reading Quoted Input	✗ Reads only first word	✓ Reads full string with spaces
Handling Escaped Quotes	✗ Manual (\")	✓ Automatic (" → \")
Handling Backslashes	✗ Manual (\\)	✓ Automatic (\ → \\)
Improves Code Readability?	✗ No	✓ Yes

20. The `std::integer_sequence`

- What Problem Does `std::integer_sequence` Solve?
 - Imagine you want to do something with a sequence of numbers, but you need to do it at **compile time**. This means the compiler has to figure out the sequence and how to use it before the program even runs. **Traditional loops and arrays won't work for this because they operate at runtime.**
 - `std::integer_sequence` provides a way to represent a sequence of integers that the compiler can work with. It's a **template class**, meaning you can specify the type of integer (usually `int`) and the sequence itself.
 - It is part of the **type-traits** library (`<utility>` header) and is mainly used in **template metaprogramming**.
- How `std::integer_sequence` Works?
 - Think of `std::integer_sequence` as a "bag" of integers known at compile time. You don't access the elements individually like you would with an array (e.g., `my_sequence[2]`). Instead, you use template metaprogramming techniques to work with the entire sequence.

- Example: Let's say you want to generate code that prints the numbers 0, 1, 2, and 3.
Here's how you can do it with `std::integer_sequence`:

```
// Function to print numbers
template <std::size_t... Indices>
void printNumbers(std::integer_sequence<std::size_t, Indices...>) {
    using expand = int[]; // Helper for parameter pack expansion
    (void)expand {
        0, (std::cout << Indices << " ", 0)...
    }; // Expand pack
    std::cout << '\n';
}

int main() {
    // Generate sequence 0, 1, 2, 3
    printNumbers(std::make_integer_sequence<std::size_t, 4>{});
    return 0;
}
```